

Laboratory A

Group 4 Predrag Bozhovikj, Erik Sievers

Assignment 1

Benchmarks on modified function `pJackknife m f` with injected mapping function `m` modified to be parallelised. Benchmarking is run with params `+RTS -A300m -lf -N4 -s`.

Benchmark results for the sequential map.

1	benchmarking map	(sequential)
2	time	428.7 ms (422.7 ms .. 434.4 ms)
3		1.000 R ² (1.000 R ² .. 1.000 R ²)
4	mean	433.8 ms (431.3 ms .. 435.5 ms)
5	std dev	2.931 ms (2.199 ms .. 3.280 ms)
6	variance introduced by outliers: 19% (moderately inflated)	

pmap: `par, pseq`

1	benchmarking pmap	(par, pseq)
2	time	188.2 ms (168.4 ms .. 215.6 ms)
3		0.987 R ² (0.969 R ² .. 1.000 R ²)
4	mean	190.4 ms (184.4 ms .. 198.1 ms)
5	std dev	9.223 ms (4.376 ms .. 13.31 ms)
6	variance introduced by outliers: 14% (moderately inflated)	

rpmmap: `rpar, rseq`

1	benchmarking rpmmap	(rpar, rseq)
2	time	179.6 ms (176.6 ms .. 183.5 ms)
3		1.000 R ² (0.999 R ² .. 1.000 R ²)
4	mean	177.4 ms (176.1 ms .. 178.5 ms)
5	std dev	1.797 ms (1.348 ms .. 2.299 ms)
6	variance introduced by outliers: 12% (moderately inflated)	

chunkMap: **Strategy**

Run with strategy `parListChunk 100 rdeepseq`.

1	benchmarking chunkMap	(strategies)
2	time	183.3 ms (173.5 ms .. 192.7 ms)
3		0.997 R ² (0.989 R ² .. 1.000 R ²)
4	mean	183.7 ms (180.4 ms .. 191.0 ms)
5	std dev	6.352 ms (1.962 ms .. 9.754 ms)
6	variance introduced by outliers: 14% (moderately inflated)	

Par monad

The implementation of map using parallel monad is in the function `_parMap` which uses a threshold parameter `t` to divide the list into `t` sublists, returning an `IVar` for each sublist.

1	benchmarking _parMap	(Par monad)
2	time	185.5 ms (179.4 ms .. 191.1 ms)
3		0.999 R ² (0.999 R ² .. 1.000 R ²)
4	mean	181.2 ms (179.3 ms .. 183.5 ms)
5	std dev	2.826 ms (1.918 ms .. 3.889 ms)
6	variance introduced by outliers: 14% (moderately inflated)	

Assignment 2

Let's first look at the sorting “ benchmarking divide and conquer sort time 1.554 s (943.7 ms .. 2.768 s) 0.934 R² (0.892 R² .. 1.000 R²) mean 1.310 s (1.209 s .. 1.469 s) std dev 173.7 ms (48.20 ms .. 235.1 ms) variance introduced by outliers: 24% (moderately inflated)

benchmarking built-in sort time 990.5 ms (753.3 ms .. 1.277 s) 0.990 R² (0.968 R² .. 1.000 R²) mean 1.013 s (960.6 ms .. 1.071 s) std dev 70.11 ms (23.90 ms .. 95.13 ms) variance introduced by outliers: 20% (moderately inflated) ”

As we can see, the divide and conquer sort is a bit slower than the standard implementation.

Let's look at the search

1	benchmarking divide and conquer search		
2	time	183.0 ms	(140.8 ms .. 214.0 ms)
3		0.980 R ²	(0.957 R ² .. 1.000 R ²)
4	mean	188.2 ms	(176.7 ms .. 205.3 ms)
5	std dev	17.31 ms	(9.966 ms .. 27.02 ms)
6	variance introduced by outliers: 16% (moderately inflated)		
7			
8	benchmarking built-in search		
9	time	8.108 ms	(7.600 ms .. 8.991 ms)
10		0.887 R ²	(0.724 R ² .. 0.989 R ²)
11	mean	8.783 ms	(8.419 ms .. 9.750 ms)
12	std dev	1.599 ms	(657.9 μs .. 2.855 ms)
13	variance introduced by outliers: 80% (severely inflated)		

The search is a lot slower than the default implementation. It may be a good idea to not spark everything, but instead chunk it after a given depth to improve performance.


Assignment 3

3.1 parBuffer


`parBuffer` takes an integer and a strategy as a parameter and then puts that number of sparks in a buffer. Once a spark is done, another can be added.

3.2 Compare parListChunk with parBuffer

Let's start by looking at how `parListChunk` creates and consumes sparks from the spark pool




`parListChunk` creates all sparks right away




`parListChunk` then consumes the sparks from the spark pool

Now, let's contrast that with `parBuffer`




`parBuffer`, on the other hand, creates a number of sparks and then creates new sparks as the created ones finish.




`parBuffer` as we can see, the number of sparks is somewhat stable. As sparks are consumed, new sparks are added to the spark pool

3.3 Combine chunking with `parBuffer`

Let's see how our `parBufferChunk` performed





It's similar to `parBuffer` in that it has a somewhat stable number of sparks throughout the whole execution. However, with the parameters we chose (buffer size 10, chunk size 100) the `parBuffer` allocates 16000 sparks, whereas the `barBufferChunk` only allocates 160 sparks. This reduced overhead makes `barBufferChunk` almost twice as fast as `parBuffer` (39s vs 17s).